

STA414 Assignment 3

Jeff Blair: 1002177057
blairjef

Collaborator: Jai Aggarwal

April 2020

1 Implementing the Model

```
a. log_prior(z) = factorized_gaussian_log_density(0, 0, z)

b. decoder = Chain(Dense(Dz, Dh, tanh), Dense(Dh, Ddata))

c. function log_likelihood(x,z)
    """ Compute log likelihood log_p(x|z) """
    theta = decoder(z)
    return sum(bernoulli_log_density(theta, x), dims=1) # return likelihood for each element in batch
end

d. joint_log_density(x, z) = log_prior(z) .+ log_likelihood(x, z)
```

2 Amortized Approximate Inference and training

```
a. encoder = Chain(Dense(Ddata, Dh, tanh), Dense(Dh, 2*Dz), unpack_gaussian_params)

b. log_q(q_mu, q_logsigma, z) = factorized_gaussian_log_density(q_mu, q_logsigma, z)

c. function elbo(x)
    q_mu, q_logsigma = encoder(x)
    z = sample_diag_gaussian(q_mu, q_logsigma)
    joint_ll = joint_log_density(x, z)
    log_q_z = log_q(q_mu, q_logsigma, z)
    elbo_estimate = mean(joint_ll .- log_q_z)
    return elbo_estimate
end

d. function loss(x)
    return -elbo(x) #TODO: scalar value for the variational loss over elements in the batch
end
```

```

e. function train_model_params!(loss, encoder, decoder, train_x, test_x; nepochs=10)
    # model params
    ps = Flux.params(encoder, decoder)
    # ADAM optimizer with default parameters
    opt = ADAM()
    # over batches of the data
    for i in 1:nepochs
        for d in batch_x(train_x)
            # compute gradients with respect to variational loss over batch
            gs = Flux.gradient(ps) do
                return loss(d)
            end

            Flux.Optimise.update!(opt, ps, gs)
        end
        if i%1 == 0 # change 1 to higher number to compute and print less frequently
            @info "Test loss at epoch $i: $(loss(batch_x(test_x)[1]))"
        end
    end
    @info "Parameters of encoder and decoder trained!"
end

## Train the model
train_model_params!(loss,encoder,decoder,train_x,test_x, nepochs=100)
Info: Test loss at epoch 100: 155.76130769378668
@ Main In[144]:20
Info: Parameters of encoder and decoder trained!
@ Main In[144]:23


```

3 Visualizing Posteriors and Exploring the Model

```

a. plot_list = Any[]
for i in 1:10
    z = sample_diag_gaussian([0, 0],0)
    # Use generative model to compute the bernoulli means over the pixels of x given z.
    theta = decoder(z)
    means = exp.(theta)./(1 .+ exp.(theta))
    # Plot as a greyscale image
    push!(plot_list, plot(mnist_img(means)))
    # Sample a binary image x from this product of Bernoullis. Plot this sample
    out = sample_bernoulli.(means)
    push!(plot_list, plot(mnist_img(out), size=(4000, 800)))
end
display(plot(plot_list..., layout=grid(2,10)))

```

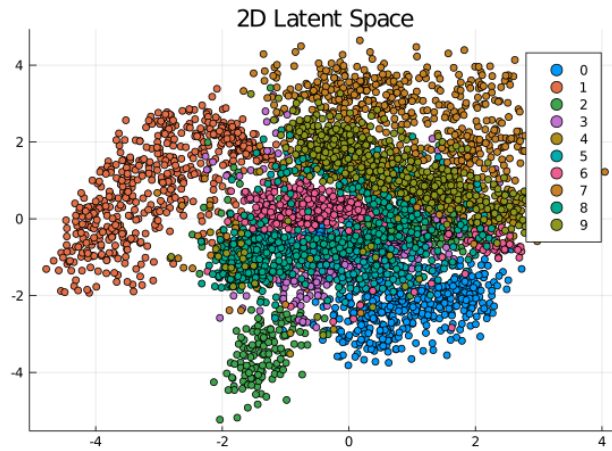


(a) Means vs Binary

```

b. encodings = encoder(train_x)
   means = encodings[1]
   scatter(means[1,:], means[2,:], group=train_label, title = "2D Latent Space")

```



(a) 2D latent space

```


c. function li(za, zb, alpha=0.5)
    return (alpha .* za) .+ ((1 .- alpha) .* zb)
end

labels = []
for i in 1:3
    push!(labels, sample([0,1,2,3,4,5,6,7,8,9], 2; replace=false))
end

# 3 pairs of images
images = []
for pair in labels
    s1 = train_x[:, train_label .== pair[1]]
    s2 = train_x[:, train_label .== pair[2]]
    image1 = s1[:, rand(1:size(s1)[2])]
    image2 = s2[:, rand(1:size(s2)[2])]
    push!(images, [encoder(image1)[1], encoder(image2)[1]])
end

# 10 interpolations per pair
plot_list = []
for pair in images
    for alpha in 0.1:0.1:1
        interp = li(pair[2], pair[1], alpha)
        theta = decoder(interp[:,1])
        means = exp.(theta)./(1 .+ exp.(theta))
        # Plot as a greyscale image
        push!(plot_list, plot(mnist_img(means), size=(4000,800)))
    end
end
print(labels)
display(plot(plot_list..., layout=grid(3,10)))

```



(a) Linear Interpolation

Labels for the first row were $8 \rightarrow 7$, second row $6 \rightarrow 2$, third row $3 \rightarrow 1$

4 Predicting the Bottom of Images given the Top

```
a. # Write a function that computes p(z/top half of the image)
function top_half(x)
    y = reshape(x, 28, 28,:)
    z = y[:,1:14,:]
    return reshape(z, (392,:))
end

# Write a function that computes log p(top half of image x | z)
function top_log_px_z(x, z)
    theta = decoder(z)
    return sum(bernoulli_log_density(top_half(theta), top_half(x)), dims=1)
end

# Combine this likelihood with the prior to get a function that computes the joint log density
function joint_log_density_top(x, z)
    return log_prior(z) + top_log_px_z(x, z)
end

b. # Initialize mu and sigma
mu = randn(2)
logsigma = randn(2)

# Write a function that computes ELBO over K samples
function elbo_th(x, opts, k)
    mu, logsigma = opts
    z = sample_diag_gaussian(repeat(mu, 1, k), logsigma)
    joint_ll = joint_log_density_top(x, z)
    log_q_z = log_q(mu, logsigma, z)
    elbo_estimate = mean(joint_ll .- log_q_z)
    return elbo_estimate
end

function loss_th(x, opts, k)
    return -elbo_th(x, opts, k)
end

# Training
s = train_x[:, rand(1:1000)]
function train(opts, t_x, iters=200, lr=1e-2, k=10)
    opts_cur = opts
    # x = t_x[:, rand(1:1000)]
    x = t_x
```


```

for i in 1:iters
    opts_grad = gradient(opts_cur -> loss_th(x, opts_cur, k), opts_cur)
    opts_cur[1] -= lr .* opts_grad[1][1]
    opts_cur[2] -= lr .* opts_grad[1][2]
    @info "Elbo:" elbo_th = elbo_th(x, opts_cur, k)
end
return opts_cur
end

# Train Paramters
opts = mu, logsigma
opts = train(opts, s)

# Plotting Isocontours
mu, logsigma = opts
joint(z) = exp(joint_log_density_top(s, z))
post(z) = exp(log_q(mu, logsigma, z))
plot(title="Joint Density vs Approx Posterior Contours")
skillcontour!(joint)
skillcontour!(post)

```



(a) Contours

```

z = sample_diag_gaussian(mu, logsigma)

# Use generative model to compute the bernoulli means over the pixels of x given z.
theta = decoder(z)
means = exp.(theta)./(1 .+ exp.(theta))


```

```

# Concatenate top half and predicted bottom half
top = top_half(s)
bot = reshape(means, 28, 28, :)
bot = bot[:,15:end,:]
bot = reshape(bot, (392,:))

img = vcat(top, bot)
plot_list = [plot(mnist_img(img[:,1])), plot(mnist_img(s))]
display(plot(plot_list...))
savefig("4be")

```



(a) Predicted bottom half

- c. a) True (Yes)
 b) False (No)
 c) True (Yes)
 d) False (No)
 e) True (Yes)